

Safe Agentic AI Framework and Ecosystem Roadmapping (SAFER) Workshop to create a Trustworthy Agentic AI Roadmap

IJCAI-ECAI 2026 Workshop

August 15-17, 2026

Bremen, Germany

Call for Submissions and Participation

As AI models increase in capability, agency and autonomy, their impact on individuals and society grows. This interdisciplinary workshop aims to attract researchers, practitioners, policy makers and students from diverse fields including computer science, sociology, philosophy, and law to collaboratively create a roadmap considering the scientific and technical, security, economic, environmental, regulatory, social and ethical implications of Agentic AI to advance responsible if not beneficent agents and their enabling ecosystem. This interactive workshop will include a mix of invited and submitted papers and is designed to foster the interdisciplinary dialogue that is highly valued at IJCAI-ECAI and essential to advance progress responsibly. This global topic is timely, interdisciplinary, and necessary, addressing a critical need in the AI community and society at large.

Submissions:

We invite submissions that include surveys of existing agent gaps or forecasts, late breaking insights into future roadmaps and speculative but clearly articulated agent futures. Submissions will be evaluated based on the IJCAI-ECAI criteria of novelty, originality, technical correctness, clarity, significance of results, and potential impact. Given the roadmap focus of the workshop, we will encourage results that both carefully articulated gaps, evidence-based forecasts as well as visionary but clearly articulated desired future capabilities or outcomes. Using the standard IJCAI format and templates (available [here](#)), **submissions should be made to chairingtool.com/conferences/safer2026/main-track?role=author.** The workshop web site is github.com/mtmaybury/SafeAgents as either a short paper (2-4 pages) or long paper (no longer than 7 pages plus 2 pages for references). Submissions do not need to be anonymized.

Deadlines:

Workshop Paper Submission: May 4, 2026 (IJCAI-ECAI Main Conf Paper decisions April 29, 2026)

Workshop Paper Notification: June 15, 2026

Camera Ready Copy: July 13, 2026

Workshop (1 of these days): Saturday/Sunday/Monday, August 15-17, 2026.

Organizing Committee

- Dr. [Mark Maybury](#), VP Commercialization, Lockheed Martin; Board Director, Advanced Cyber Security Center. mark.t.maybury@lmco.com (chair and primary contact)
- Prof. [Wolfgang Wahlster](#), Founding Director and CEA, DFKI. Wolfgang.Wahlster@dfki.de
- Ms. [Josephine Liu](#), Chair, APAC Tech Policy Committee, Association for Computing Machinery (ACM), yuyinl@acm.org
- Dr. [Francesca Rossi](#), IBM Fellow, IBM Global Leader for Responsible AI and AI Governance, T.J. Watson Research Center, Yorktown Heights, USA Francesca.Rossi2@ibm.com

- Dr. [Mihai Christodorescu](#), Research Scientist, Security and Privacy, Google, christodorescu@google.com

This highly interactive and interdisciplinary workshop will provide a unique forum to jointly and iteratively refine a roadmap which will be publicly shared to address these complex challenges.

Goal and Product: Trustworthy Agentic AI Roadmap (TAP) articulating, current *critical gaps* or *challenges*, planned key *capabilities* and desired *outcomes*.

Technical Book: If sufficient interest from participants, a collection of the best contributions will be edited into a peer reviewed technical book.

Critical Questions of Interest: These include but are not limited to:

- **Agentic AI:** Agentic AI, or goal-seeking AI systems that can act autonomously to achieve specific objectives are expanding in capability and application. While agents promise to revolutionize everything from education to customer service to healthcare, they introduce new challenges such as control (e.g., delegation of authority, mixed-initiative), accountability, and verification and validation of evolving learning systems as well as perils arising from bias, brittleness and belligerence [Maybury 2025a,b] and resultant harms (e.g., financial, physical, social, reputational). How do we ensure and assure progress and avoid unintended and undesirable consequences?
- **Models and Autonomy:** The global deployment and use of large language/process models, small models for edge devices, and the introduction of large world models is accelerating new cases and advancing agents and autonomous systems across the enterprise and at the edge. Their computational and energy appetite for model training and inference drive a need for resource-efficient infrastructure and processing. How can we accelerate the state of the art in algorithms, architectures, interoperability (e.g., Model Context Protocol) and security to simultaneously advance performance, usability, prosperity, transparency, and safety?
- **Society, Safety and Sustainability.** Current LLMs suffer from being biased, brittle, and baroque with risk of significant harms. A rapidly expanding number of agents and humanoid robots add new risks of belligerence. Beginning to address these concerns, societies have introduced national guidelines (e.g., NIST AI RMF), regulations (e.g., DoD 3000.09, Chinese Regulation on the Governance of Generative AI), strategies (e.g., Germany's AI Action Plan, US AI Action Plan, UK AI Strategy), laws (e.g., EU AI Act, GDPR), and sector-specific safety and security guidelines. What are effective methods for responsible development, deployment and operation that can mitigate risks to society? How can governance frameworks, standards, and regulatory mechanisms be effectively implemented and coordinated across jurisdictions and institutions to ensure responsible development and deployment of agentic AI systems?

Workshop Themes

Based on submissions and final participants, the workshop will be structured around several key sub-themes of agent safety to ensure a comprehensive and focused discussion including:

- **Ethical Foundations and Governance:** theoretical and philosophical underpinnings of ethical AI for agents. Topics can include:
 - Assessing current and new ethical frameworks for autonomous and agentic systems.
 - The role of law, regulation and policy in guiding agent development.

- Philosophical debate around agent consciousness, rights and responsibilities.
- **Accuracy, Bias, Fairness, Transparency:** A core technical and social component of the workshop will address:
 - Methods for improving agent accuracy and self-awareness (e.g., limitations, confidence).
 - Successful approaches to detecting and mitigating dangers such as bias and hallucination in large datasets and models.
 - Techniques for creating more transparent and interpretable models and agents.
 - Evaluating the impact of agents on different demographic groups.
- **Human-Agent Collaboration and Trust:** As intelligent agents becomes more integrated into daily life, understanding the dynamics of human-agent collaboration is crucial, including:
 - Designing intuitive and trustworthy user interfaces for agentic AI.
 - The role of agents in decision support systems and associated risks (e.g., awareness, over/under confidence, deskilling).
 - Governance frameworks for autonomous decision-making by AI agents, including responsibility allocation, oversight mechanisms, and human-in-the-loop or human-on-the-loop models.
 - Strategies for building public trust and understanding of agent technologies.
- **Real-world Applications and Case Studies:** Case studies from various industries such as:
 - Responsible development and deployment in low to high risk use cases
 - Responsible agents in high assurance applications with human, financial, and security risks.
 - Ethical and operational challenges in using agents in high assurance and regulated industries such as healthcare, disaster management, pharmaceutical, financial services, telecommunications, logistics, energy, agriculture, defense, and space.
 - The use of AI in creative fields and its impact on human learning, training, and creativity.

Target Participants

- **Academic Researchers:** Computer scientists, ethicists, sociologists, and legal scholars.
- **Industry Practitioners:** AI engineers, product managers, and data scientists working on real-world applications.
- **Policymakers/Regulators:** Officials and orgs concerned with AI/Agent governance and safety.
- **Students:** Graduate and undergraduate students interested in social and ethical dimensions of AI.

Workshop Format and 1 Day Schedule

- Prior to the workshop, the program committee will issue a call for submissions of position papers for potential participants, requiring a position paper with recommended inputs to a Trustworthy Agentic AI Roadmap (TAP) articulating key needs, gaps, capabilities and outcomes (including time frames when these expected), and road blocks or challenges.
- The 1-day workshop is first of its kind to create safe agent roadmap and will include no more than 30-40 carefully selected participants (based on submitted position papers or proposed posters in the IJCAI format) who are experts in agents and safety to enable rich engagement. The workshop will kick off with a draft roadmap created from the participants submitted position papers, invited talks by leaders of frontier models (e.g., OpenAI, Anthropic, Google, Microsoft) and industrial domain models (e.g. SAP, Siemens, Lidl) as well as 2 or 3 invited talks (spread across the morning and afternoon) selected by the program committee and selected participant presentations organized by themes (e.g., agentic AI, models, use cases safety vulnerabilities and mitigations). If submissions support this, we will also have a short poster session in the afternoon after lunch to

maximize participant engagement. We have also proposed an IJCAI-ECAI main conference panel on new directions for safe agents mid-morning after with select members of the program committee and/or invited speakers and will include Q&A from all, orchestrated by the Workshop Chair. During the day, new gaps, opportunities, and expected outcomes will be added to the evolving roadmap. The day will conclude sharing back the final draft of the joint roadmap.