

Safe Agentic AI Framework and Ecosystem Roadmapping (SAFER) Workshop to create a [Trustworthy Agentic AI Roadmap](#)

IJCAI-ECAI 2026 Workshop AGENDA

Sunday August 16, 2026

Bremen, Germany

Register at <https://registration.ijcai.org>

8:30 Registration Opens

[8:45 Early Arrival Sneak Preview of Roadmap]

9:00 Welcome – Dr. Mark Maybury (Workshop Chair) – [Toward A Secure Agent Roadmap](#)

9:10 **Paper Session 1: *Safety and Risk*** [Papers: 15 min presentation, 5 min Q&A]

[Neuro-Symbolic Agents for Safe, Sovereign and Auditable AI Systems on the Shop Floor](#), Daniel Porta, Gerhard Sonnenberg, Gerd Herzog, Sönke Knoch, Tim Schwartz and Wolfgang Wahlster. German Research Center for Artificial Intelligence (DFKI) {firstname.lastname}@dfki.de

[Risk Perspectives in Distributed AI Systems for Complex Environments](#). André Meyer-Vitali (andre@meyer-vitali.com) and Rodica Mihai (rodica.g.mihai@norceresearch.no). German Research Center for Artificial Intelligence (DFKI)

9:50 **Paper Session 2: *Governance*** [Papers: 15 min presentation, 5 min Q&A]

[Control Before Consequence: Reframing Infrastructure and Governance for Agentic AI Systems - charlemagnevolt-hub](#) Dr. Torsten Bernasco, drtorstenbernasco@gmail.com, Aramvolt

[Adaptive Agency: A Human-AI Co-governance Framework for Secure Grassroots Governance in Complex Urban Environments](#). Xin Ma celia.ma0829@gmail.com Beijing Academy of Social Sciences

10:30-11:00 Coffee Break

11:00 **Paper Session 2: *Governance*** Governance frameworks for the safe development, deployment, and competition in emerging technologies, such as AI. Ms. [Josephine Liu](#), Chair, APAC Tech Policy Committee, Association for Computing Machinery.

11:20: **Invited Speaker:** Asago: [Open Source Community for Model Security](#). Dr. Stuart Battersby, Red Hat, *AI Safety & Model Evaluation Architect* sbatters@redhat.com

11:50: Brief Roadmap Discussion – All

12:00 LUNCH (Not provided by IJCAI) [Roadmap Discussions Encouraged]

13:00 **Paper Session 2: Governance** Mitigating Risks Across the Life Cycle of Generative Artificial Intelligence. Mark Maybury. Mark.t.maybury@lmco.com. Lockheed Martin

13:20 Review of Draft Roadmap - Everyone

14:00 **Paper Session 3 Trust Building: Proving AI is Safe and Culturally Appropriate**

[Papers: 15 min presentation, 5 min Q&A]

Cryptographic certificates of validity for trustworthy AI. Murdoch J. Gabbay, Heriot-Watt University. M.gabbay@hw.ac.uk [<https://arxiv.org/abs/2606.23768>]

Designing Identity Before Tools: Human Identity and Trust as a Missing Dimension in Agentic AI Safety. Marco Nieschka and Kordula Hildebrandt, HILDEBRANDT FILM GmbH marco.nieschka@hildebrandt-film.de and kordula.hildebrandt@hildebrandt-film.de.

14:40 Final Review of Roadmap

15:30-16:00 Coffee Break

16:00 Panel

- Dr. Mark Maybury, VP Commercialization, Lockheed Martin; Board Director, Advanced Cyber Security Center. mark.t.maybury@lmco.com (moderator)
- Prof. Wolfgang Wahlster, Founding Director and CEA, DFKI. Wolfgang.Wahlster@dfki.de
- Ms. Josephine Liu, Chair, APAC Tech Policy Committee, Association for Computing Machinery (ACM), yuyinl@acm.org
- Dr. Mihai Christodorescu, Research Scientist, Security and Privacy, Google, christodorescu@google.com
- Dr. Stuart Battersby, Red Hat, *AI Safety & Model Evaluation Architect* sbatters@redhat.com

17:00 Final Roadmap, Concluding Discussions, Future Directions

17:45 Closing

18:00 Public Panel on Safe Agents and Workshop Results - University Auditorium (capacity 500), affectionately known locally as the "Keksdose" (cookie tin). Register for the free public panel with a free ticket here: <https://2026.ijcai.org/outreach/>.

Unable to attend: Why Aligned Models Are Not Enough: Toward Channel-Level Safety for Agent-to-Agent Communication. Amr Gomaa, German Research Center for Artificial Intelligence (DFKI), amr.gomaa@dfki.de.

Dr. Francesca Rossi, IBM Fellow, IBM Global Leader for Responsible AI and AI Governance, T.J. Watson Research Center, Yorktown Heights, USA Francesca.Rossi2@ibm.com. Recent AI Governance Papers <https://francescarossiai.substack.com/p/ai-governance-is-not-done> and <https://francescarossiai.substack.com/p/trust-in-ai-is-a-question-of-governance>

Asago: RedHat – Stuart Battersby sbatters@redhat.com

- Press release: <https://www.redhat.com/en/about/press-releases/red-hat-launches-asago-community-automate-ai-safety-and-governance-policy-production>
- Blog post: <https://www.redhat.com/en/blog/introducing-asago-open-source-ai-safety-and-governance-orchestration>
- Stuart's LinkedIn post: https://www.linkedin.com/posts/stuartbattersby_super-excited-to-be-launching-asago-an-open-source-share-7490393171731042304-gopM

Dropped:

13:30pm **Invited Speaker 2:** Nicholas Carlini, Anthropic, formerly Google/DeepMind (Invited)
nicholas@carlini.com

14:40 **Invited Speaker 3:** Dr. Florian Tramèr (ETH Zurich). [AgentDojo: A Dynamic Environment to Evaluate Prompt Injection \(2025\)](#). florian.tramer@inf.ethz.ch (Invited)